

REQUESTS & LIMITS IN KUBERNETES

Requests and Limits are part of **resource management** and help control the CPU and memory (resources) consumed by containers within a pod.

1. Resource Requests

- A **request** is the amount of CPU or memory guaranteed to a container.
- Kubernetes uses requests to determine which node can schedule a pod.
- If a node has enough resources to fulfill the request, the pod will be scheduled on that node.

2. Resource Limits

- A **limit** is the maximum amount of CPU or memory a container can use.
- If a container exceeds its memory limit, it will be terminated, and if it exceeds its CPU limit, its usage will be throttled.

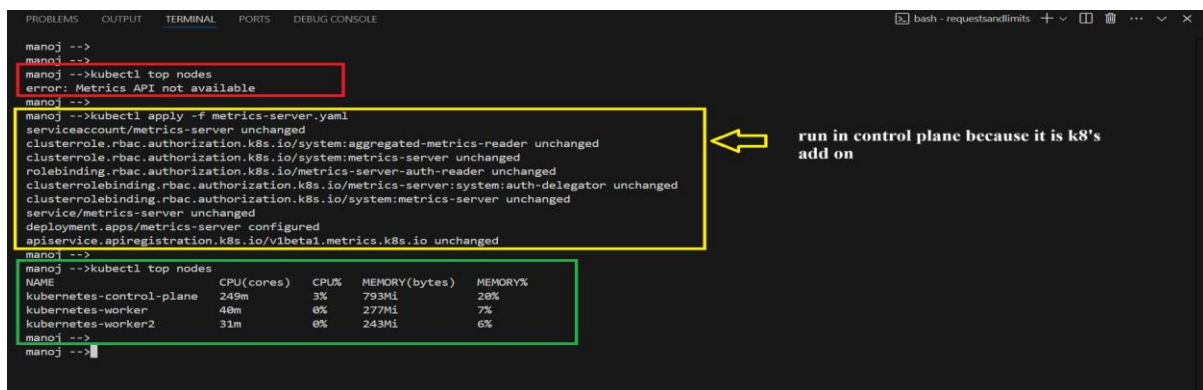
→ Some of the issues we can face in request and limits are:

- **Insufficient Memory (or Insufficient CPU):** The node doesn't have enough memory or CPU available to fulfill the pod's resource request.
- **OOMKilled:** The container was terminated because it exceeded its memory limit.
- **PodFitsResources:** The pod can't be scheduled because no node has sufficient resources to meet the pod's requests.
- **ResourceQuotaExceeded:** The pod exceeds the resource quota set for the namespace, preventing it from being created.
- **LimitRangeExceeded:** The pod's resource requests or limits exceed the constraints defined by the LimitRange policy in the namespace.
- **PodEvicted:** The pod was evicted due to resource pressure on the node, often because of insufficient memory.
- **ContainerCannotRun:** The container fails to start due to insufficient resources or incorrect configuration.
- **FailedScheduling:** Kubernetes was unable to find a suitable node to schedule the pod due to resource constraints.

For metric-server yml use the link below to find the code

https://drive.google.com/file/d/1CIbZkM_71xn4psUrGeIhfWTaKfD_5oZu/view?usp=drive_link

Run the metric-server on kube-system namespace present in the control plane



```
manoj -->
manoj -->
manoj -->kubectl top nodes
error: Metrics API not available
manoj -->
manoj -->kubectl apply -f metrics-server.yaml
serviceaccount/metrics-server unchanged
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader unchanged
clusterrole.rbac.authorization.k8s.io/system:metrics-server unchanged
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader unchanged
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator unchanged
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server unchanged
service/metrics-server unchanged
deployment.apps/metrics-server configured
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io unchanged
manoj -->
manoj -->kubectl top nodes
NAME                CPU(cores)   CPU%   MEMORY(bytes)  MEMORY%
kubernetes-control-plane 249m        3%    793Mi          28%
kubernetes-worker       40m         0%    277Mi          7%
kubernetes-worker2      31m         0%    243Mi          6%
manoj -->
manoj -->
```

run in control plane because it is k8's add on

REQUESTS & LIMITS IN KUBERNETES

Case 1: Request memory “100Mi” and limit the memory to “200Mi”. with in the limit POD is created.

```
! mem-1.yaml X
requestsandlimits > ! mem-1.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: memory-demo
5    namespace: memory-example
6  spec:
7    containers:
8      - name: memory-demo-ctr
9        image: polinux/stress
10       resources:
11         requests:
12           memory: "100Mi"
13         limits:
14           memory: "200Mi"
15       command: ["stress"]
16       args: ["--vm", "1", "--vm-bytes", "150M", "--vm-hang", "1"]
```

this the limit of the resource 200mi

with in the limit POD is creating

```
PROBLEMS  OUTPUT  TERMINAL  PORTS  DEBUG CONSOLE
bash - requestsandlimits + - + ... x

manoj -->
manoj -->
manoj --># now let do stress testing
manoj -->
manoj -->kubectl create ns memory-example
Error from server (AlreadyExists): namespaces "memory-example" already exists
manoj -->
manoj -->kubectl apply -f mem-1.yaml
pod/memory-demo created
manoj -->
manoj -->kubectl get pod -n memory-example
NAME      READY   STATUS    RESTARTS   AGE
memory-demo 1/1     Running   0           19s
manoj -->
manoj -->kubectl top pod memory-demo -n memory-example
NAME      CPU(cores)   MEMORY(bytes)
memory-demo 66m          150Mi
manoj -->
manoj -->kubectl top pod -n memory-example
NAME      CPU(cores)   MEMORY(bytes)
memory-demo 67m          150Mi
manoj -->
manoj -->
```

POD is running successfully with in the limit of the resources

Case 2: Request memory “50Mi” and limit the memory to “100Mi”. Although the container tries to allocate 250 MiB of memory, Kubernetes enforces a 100 MiB memory limit, so the container will likely be terminated due to exceeding the limit.

```
! mem-2.yaml X
requestsandlimits > ! mem-2.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: memory-demo-2
5    namespace: memory-example
6  spec:
7    containers:
8      - name: memory-demo-2-ctr
9        image: polinux/stress
10       resources:
11         requests:
12           memory: "50Mi"
13         limits:
14           memory: "100Mi"
15       command: ["stress"]
16       args: ["--vm", "1", "--vm-bytes", "250M", "--vm-hang", "1"]
```

here the limit is 100 mb

Out of Memory that is we are using more then the memory

REQUESTS & LIMITS IN KUBERNETES

```
manoj -->
manoj -->
manoj -->kubectl get pod -n memory-example
No resources found in memory-example namespace.
manoj -->
manoj -->
manoj -->kubectl apply -f mem-2.yaml
pod/memory-demo-2 created
manoj -->
manoj -->
manoj -->kubectl get pod -n memory-example
NAME          READY   STATUS    RESTARTS   AGE
memory-demo-2  0/1     Error     0           4s
manoj -->
```

POD is the ERROR state

```
manoj -->
manoj -->
manoj -->kubectl get pod -n memory-example
NAME          READY   STATUS             RESTARTS   AGE
memory-demo-2  0/1     CrashLoopBackOff   7 (52s ago) 12m
manoj -->
manoj -->
```

```
--vm-bytes
250M
--vm-hang
State:          Waiting
Reason:         CrashLoopBackOff
Last State:     Terminated
Reason:         Error
Exit Code:      1
Started:        Tue, 08 Oct 2024 21:43:07 +0530
Finished:       Tue, 08 Oct 2024 21:43:07 +0530
Ready:          False
Restart Count:  6
Limits:
  memory: 100Mi
Requests:
  memory: 50Mi
Environment:    <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-htkd8 (ro)
Conditions:
  Type                               Status
  PodReadyToStartContainers         True
  Initialized                       True
  Ready                             False
  ContainersReady                   False
  PodScheduled                      True
Volumes:
  kube-api-access-htkd8:
    Type:              Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
```

If the container exceeds the **100 MiB memory limit**, it will be terminated by Kubernetes (likely with an OOMKilled error).

```
ContainersReady      False
PodScheduled         True
Volumes:
  kube-api-access-htkd8:
    Type:              Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:      kube-root-ca.crt
    ConfigMapOptional:  <nil>
    DownwardAPI:        true
QoS Class:           Burstable
Node-Selectors:      <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                     node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type     Reason      Age    From          Message
  ----     -
  Normal   Scheduled   10m    default-scheduler   Successfully assigned memory-example/memory-demo-2 to kubernetes-worker2
  Normal   Pulled      10m    kubelet         Successfully pulled image "polinux/stress" in 2.277s (2.277s including waiting). Image size: 4841495
  bytes.
  Normal   Pulled      10m    kubelet         Successfully pulled image "polinux/stress" in 2.042s (2.042s including waiting). Image size: 4841495
  bytes.
  Normal   Pulled      9m53s  kubelet         Successfully pulled image "polinux/stress" in 2.308s (2.308s including waiting). Image size: 4841495
  bytes.
  Normal   Created     9m21s  (x4 over 10m)    kubelet         Created container memory-demo-2-ctr
  Normal   Started     9m21s  (x4 over 10m)    kubelet         Started container memory-demo-2-ctr
  Normal   Pulled      9m21s  kubelet         Successfully pulled image "polinux/stress" in 2.239s (2.239s including waiting). Image size: 4841495
  bytes.
  Normal   Pulling     8m28s  (x5 over 10m)    kubelet         Pulling image "polinux/stress"
  Warning  BackOff     9s (x48 over 10m) kubelet         Back-off restarting failed container memory-demo-2-ctr in pod memory-demo-2_memory-example(84b38b8c-7
  s86-4260-8eb9-f37a79021f07)
manoj -->
manoj -->
```

REQUESTS & LIMITS IN KUBERNETES

Case 3: Request and Limit resources are specified here are more then the resource[memory] present in node

```
! mem-3yaml X
requestsandlimits > ! mem-3yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: memory-demo-3
5    namespace: memory-example
6  spec:
7    containers:
8      - name: memory-demo-3-ctr
9        image: polinux/stress
10       resources:
11         requests:
12           memory: "1000Gi"
13         limits:
14           memory: "1000Gi"
15         command: ["stress"]
16         args: ["--vm", "1", "--vm-bytes", "150M", "--vm-hang", "1"]
```

```
manoj -->
manoj -->
manoj -->kubectl get pod -n memory-example
No resources found in memory-example namespace.
manoj -->
manoj -->kubectl apply -f mem-3.yaml
pod/memory-demo-3 created
manoj -->
manoj -->kubectl get pod -n memory-example
NAME          READY   STATUS    RESTARTS   AGE
memory-demo-3  0/1     Pending   0           3s
```

POD is in pending state

In most clusters, such a pod would likely fail to be scheduled unless a node with very high memory capacity exists.

```

PROBLEMS OUTPUT TERMINAL PORTS DEBUG CONSOLE
--vm-bytes
150M
--vm-hang
1
Limits:
  memory: 1000Gi
Requests:
  memory: 1000Gi
Environment: <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-gw4mp (ro)
Conditions:
  Type          Status
PodScheduled   False
Volumes:
  kube-api-access-gw4mp:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:  kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:    true
QoS Class:       Burstable
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                  node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type          Reason              Age   From          Message
  ----          -
Warning        FailedScheduling    2m40s default-scheduler  0/3 nodes are available: 1 node(s) had untolerated taint {node-role.kubernetes.io/control-plane: }
; 2 Insufficient memory. Preemption: 0/3 nodes are available: 1 Preemption is not helpful for scheduling, 2 No preemption victims found for incoming pod.
manoj ->
manoj ->

```

Key Points:

- **Requests:** Minimum resources guaranteed for a container.
- **Limits:** Maximum resources a container is allowed to consume.
- If you only set a **request**, Kubernetes will not limit resource usage beyond that.
- If you only set a **limit**, the pod might not be scheduled if the node doesn't have sufficient resources.