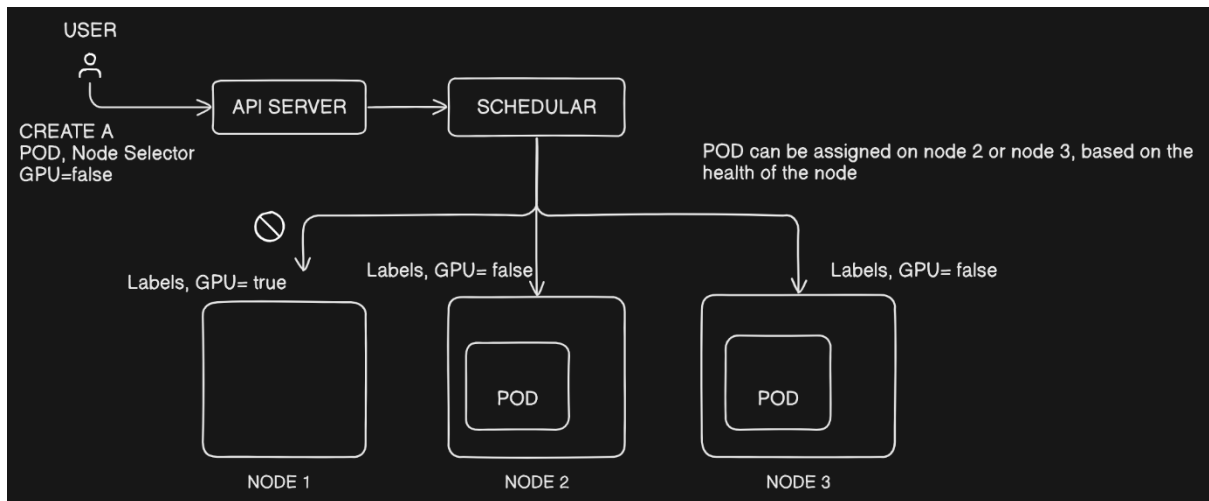


Node Selector is a simple way to constrain which nodes a pod can be scheduled on based on node labels. It allows you to ensure that a pod is only scheduled on nodes that match specific criteria.

How Node Selector Works

- **Node labels** are key-value pairs attached to nodes. You can add labels to nodes to categorize them (e.g., by hardware, software, or environment).
- A **Node Selector** in a pod's specification matches the node labels. The pod will only be scheduled on nodes that have the specified label(s).



```

! nodesselector.yaml x
tolerationandaffinityandselector > ! nodesselector.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    labels:
5      run: nginx
6    name: nginx
7  spec:
8    containers:
9      - image: nginx
10     name: nginx
11
12     nodeSelector:
13       gpu: "false"
  
```

POD will be scheduled on the matching label of the node, where 'gpu=false'

```

manoj -->kubectl get nodes
NAME                                STATUS  ROLES    AGE   VERSION
kubernetes-control-plane            Ready   control-plane  17d   v1.31.0
kubernetes-worker                   Ready   <none>        17d   v1.31.0
kubernetes-worker2                  Ready   <none>        17d   v1.31.0
manoj -->
manoj -->kubectl describe node kubernetes-worker
Name:                               kubernetes-worker
Roles:                               <none>
Labels:                               beta.kubernetes.io/arch=amd64
                                      beta.kubernetes.io/os=linux
                                      gpu=false
                                      kubernetes.io/arch=amd64
                                      kubernetes.io/hostname=kubernetes-worker
                                      kubernetes.io/os=linux
Annotations:                         kubeadm.alpha.kubernetes.io/cri-socket: unix:///run/containerd/containerd.sock
                                      node.alpha.kubernetes.io/ttl: 0
                                      volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp:                   Fri, 20 Sep 2024 22:42:18 +0530
Taints:                               <none>
Unschedulable:                       false
  
```

NODE SELECTOR AND NODE AFFINITY

```
manoj -->
manoj -->
manoj -->
manoj -->kubectl get nodes
NAME                                STATUS    ROLES    AGE    VERSION
kubernetes-control-plane            Ready    control-plane    17d    v1.31.0
kubernetes-worker                   Ready    <none>        17d    v1.31.0
kubernetes-worker2                  Ready    <none>        17d    v1.31.0
manoj -->
manoj -->kubectl apply -f nodeselector.yaml
pod/nginx created
manoj -->
manoj -->kubectl get po
NAME    READY    STATUS    RESTARTS    AGE
nginx   0/1      Pending   0            10s
manoj -->
manoj -->kubectl get po -o wide
NAME    READY    STATUS    RESTARTS    AGE    IP        NODE    NOMINATED NODE    READINESS GATES
nginx   0/1      Pending   0            22s    <none>    <none>    <none>            <none>
manoj -->
manoj -->
```

we can see still the POD is in pending state

```
manoj -->kubectl describe po nginx
Name:      nginx
Namespace: default
Priority:   0
Service Account: default
Node:      <none>
Labels:    run-nginx
Annotations: <none>
Status:    Pending
IPs:       <none>
Containers:
  Image:    nginx
  Port:     <none>
  Host Port: <none>
  Environment: <none>
  Mounts:
    /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-fsn8j (ro)
Conditions:
  Type             Status
  PodScheduled     False
Volumes:
  kube-api-access-fsn8j:
    Type:              Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:       kube-root-ca.crt
    ConfigMapOptional:   <nil>
    DownwardAPI:         true
QoS Class:           BestEffort
Node-Selectors:      gpu=false
Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                     node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type     Reason              Age           From          Message
  ----     -
Warning   FailedScheduling    2m25s        default-scheduler  0/3 nodes are available: 1 node(s) had untolered taint (node-role.kubernetes.io/control-plane: ); 2 node(s) didn't match Pod's node affinity/selector; preemption: 0/3 nodes are available: 3 Preemption is not helpful for scheduling.
manoj -->
```

```
manoj -->
manoj -->
manoj -->
manoj -->kubectl get nodes
NAME                                STATUS    ROLES    AGE    VERSION
kubernetes-control-plane            Ready    control-plane    17d    v1.31.0
kubernetes-worker                   Ready    <none>        17d    v1.31.0
kubernetes-worker2                  Ready    <none>        17d    v1.31.0
manoj -->
manoj -->kubectl get nodes --show-labels
NAME                                STATUS    ROLES    AGE    VERSION    LABELS
kubernetes-control-plane            Ready    control-plane    17d    v1.31.0    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubernetes-control-plane
kubernetes-worker                   Ready    <none>        17d    v1.31.0    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubernetes-worker,kubernetes.io/os=linux
kubernetes-worker2                  Ready    <none>        17d    v1.31.0    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubernetes-worker2,kubernetes.io/os=linux
manoj -->
manoj -->kubectl label node kubernetes-worker gpu=false
node/kubernetes-worker labeled
manoj -->
manoj -->kubectl get nodes --show-labels
NAME                                STATUS    ROLES    AGE    VERSION    LABELS
kubernetes-control-plane            Ready    control-plane    17d    v1.31.0    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubernetes-control-plane
kubernetes-worker                   Ready    <none>        17d    v1.31.0    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,gpu=false,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubernetes-worker,kubernetes.io/os=linux
kubernetes-worker2                  Ready    <none>        17d    v1.31.0    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubernetes-worker2,kubernetes.io/os=linux
manoj -->
manoj -->kubectl get po -o wide
NAME    READY    STATUS    RESTARTS    AGE    IP        NODE    NOMINATED NODE    READINESS GATES
nginx   1/1      Running   0            5m38s    10.244.1.11    kubernetes-worker    <none>            <none>
manoj -->
manoj -->
```

no labels is present on the node, that is why POD didn't scheduled on any node

adding label on the node

now we can see the label on the node

now POD got assigned onto the node, which matches the label

NODE SELECTOR AND NODE AFFINITY

```
IP: 10.244.1.11
IPs: 10.244.1.11
Containers:
  nginx:
    Container ID: containerd://36790c3899e5ff42e2b907b3b3b759c98de3f963808f7e6e4f55bc3114a46c1
    Image: nginx
    Image ID: docker.io/library/nginx@sha256:d2eb56950b84efe34f966a2b92efb1a2ea53e7e93b94cd45a27cf3cd47fc0
    Port: <none>
    Host Port: <none>
    State: Running
    Started: Tue, 08 Oct 2024 14:53:56 +0530
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-f5n8j (ro)
Conditions:
  Type              Status
  PodReadyToStartContainers  True
  Initialized        True
  Ready              True
  ContainersReady    True
  PodScheduled       True
Volumes:
  kube-api-access-f5n8j:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: True
  QoS Class: BestEffort
  Node-Selectors: gpu=false
  Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
               node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type      Reason      Age    From      Message
  ----      -
  Warning   FailedScheduling  15m    default-scheduler  0/3 nodes are available: 1 node(s) had untolerated taint [{node-role.kubernetes.io/control-plane: }], 2 node(s) didn't match Pod's node affinity/selector. Preemption: 0/3 nodes are available: 3 Preemption is not helpful for scheduling.
  Normal    Scheduled    10m    default-scheduler  Successfully assigned default/nginx to kubernetes-worker
  Normal    Pulling     10m    kubelet           Pulling image "nginx"
  Normal    Pulled      10m    kubelet           Successfully pulled image "nginx" in 1.712s (1.712s including waiting). Image size: 72958394 bytes.
  Normal    Created     10m    kubelet           Created container nginx
  Normal    Started     10m    kubelet           Started container nginx
```

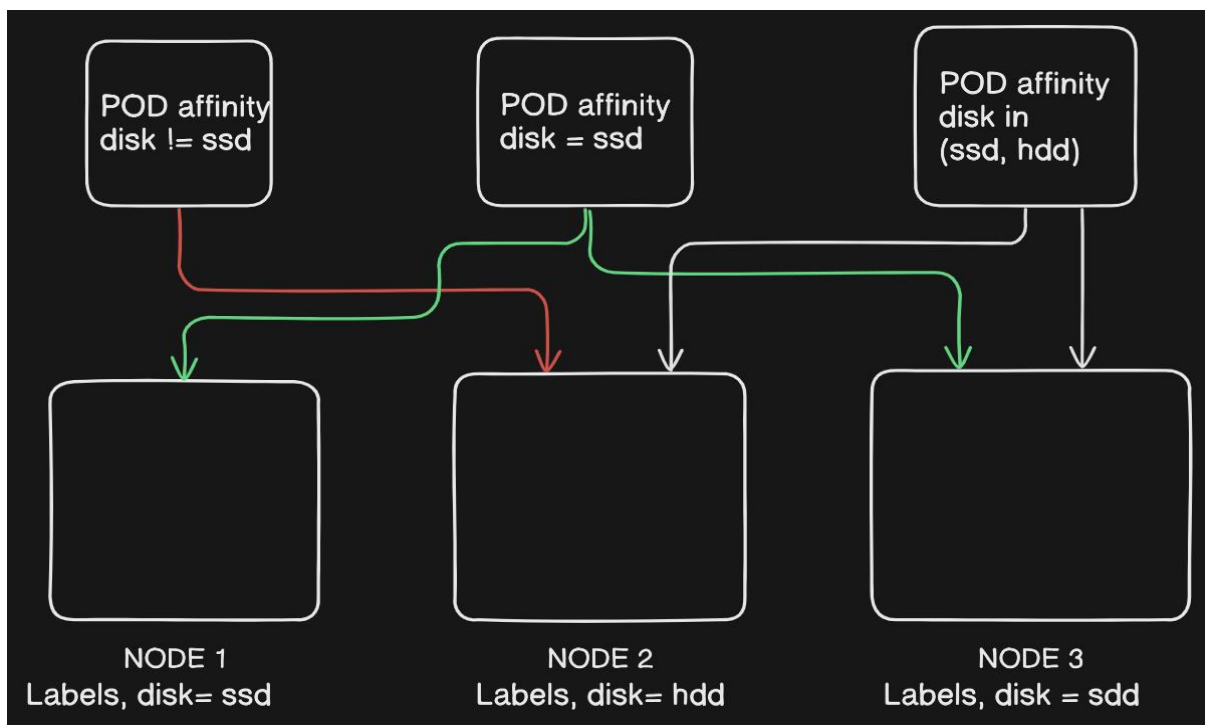
← **Node Selector**

POD got assigned onto the node

Limitations

- **Node Selector** is a simple, exact match. If you need more complex placement logic (such as multiple conditions or soft constraints), you might want to use **Node Affinity**, which offers more flexibility for scheduling decisions.

Node Affinity in Kubernetes is an advanced scheduling feature that provides more flexible rules for controlling which nodes a pod can be scheduled on, compared to the simpler **Node Selector**. It allows you to express both hard and soft constraints and supports more complex matching logic based on node labels.



NODE SELECTOR AND NODE AFFINITY

Key Features of Node Affinity

1. **Hard constraints (requiredDuringSchedulingIgnoredDuringExecution):** These are mandatory rules. If a node does not meet these conditions, the pod will not be scheduled on that node. This behavior is similar to nodeSelector, but with more complex matching options.
2. **Soft constraints (preferredDuringSchedulingIgnoredDuringExecution):** These are "preferences." Kubernetes will try to place the pod on a node that meets the soft constraint, but it's not a strict requirement. If no such nodes are available, the pod will still be scheduled on other nodes that don't meet the preference.
3. **Match Expressions:** Unlike nodeSelector, which only allows exact matches, node affinity allows you to use expressions such as:
 - **In:** Match nodes with any of the listed values for a key.
 - **NotIn:** Exclude nodes with any of the listed values for a key.
 - **Exists:** Match nodes that have the specified key, regardless of its value.
 - **DoesNotExist:** Exclude nodes that have the specified key.
 - **Gt/Lt:** Match nodes with values greater than or less than a specific number (for numerical label values).

```
! nodeaffinity.yaml x
tolerationandaffinityandselector > ! nodeaffinity.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    labels:
5      run: nginx
6    name: nginx
7  spec:
8    containers:
9      - image: nginx
10      name: nginx
11
12  affinity:
13    nodeAffinity:
14      requiredDuringSchedulingIgnoredDuringExecution:
15        nodeSelectorTerms:
16          - matchExpressions:
17            - key: disktype
18              operator: In
19              values:
20                - ssd
```

```
PROBLEMS  OUTPUT  TERMINAL  PORTS  DEBUG CONSOLE
manoj -->
manoj -->
manoj -->kubectl get po
No resources found in default namespace.
manoj -->
manoj -->kubectl get node --show-labels
NAME                                STATUS  ROLES    AGE  VERSION  LABELS
kubernetes-control-plane            Ready   control-plane  17d  v1.31.0  beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubernetes-control-plane,kubernetes.io/os=linux,node-role.kubernetes.io/control-plane,node.kubernetes.io/exclude-from-external-load-balancers
kubernetes-worker                  Ready   <none>      17d  v1.31.0  beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubernetes-worker,kubernetes.io/os=linux
kubernetes-worker2                  Ready   <none>      17d  v1.31.0  beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubernetes-worker2,kubernetes.io/os=linux
manoj -->
manoj -->
manoj -->kubectl apply -f nodeaffinity.yaml
pod/nginx created
manoj -->
manoj -->kubectl get pod -o wide
NAME    READY  STATUS   RESTARTS  AGE  IP        NODE    NOMINATED NODE  READINESS GATES
nginx   0/1    Pending  0          9s   <none>    <none>   <none>           <none>
manoj -->
manoj -->kubectl get pod -o wide
NAME    READY  STATUS   RESTARTS  AGE  IP        NODE    NOMINATED NODE  READINESS GATES
nginx   0/1    Pending  0         25s   <none>    <none>   <none>           <none>
manoj -->
manoj -->
```

↑ i don't see the node label

← POD is in still pending state

NODE SELECTOR AND NODE AFFINITY

```

[roo] --kubectl describe pod nginx
Name:          nginx
Namespace:     default
Priority:       0
ServiceAccount: default
Node:          <none>
Labels:        run/nginx
Annotations:    <none>
Status:        Pending
IP:            <none>
Containers:
  nginx:
    Image:          nginx
    Port:           <none>
    Host Port:      <none>
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-7qhv (ro)
Conditions:
  Type             Status
PodScheduled       False
Volumes:
  kube-api-access-7qhv:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:    kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:      true
QoS Class:          BestEffort
Node-Selectors:      <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                     node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type      Reason      Age    From      Message
  --      -
Warning    FailedScheduling 6m49s default-scheduler 0/3 nodes are available: 1 node(s) had untolerated taint (node-role.kubernetes.io/control-plane: 2 node(s) didn't match Pod's node affinity/selector, Preemption: 0/3 nodes are not helpful for scheduling.
Warning    FailedScheduling 92s    default-scheduler 0/3 nodes are available: 1 node(s) had untolerated taint (node-role.kubernetes.io/control-plane: 2 node(s) didn't match Pod's node affinity/selector, Preemption: 0/3 nodes are not helpful for scheduling.
[roo] -->

```

```

manoj -->
manoj -->
manoj -->kubectl get nodes

NAME                                STATUS    ROLES    AGE   VERSION
kubernetes-control-plane           Ready    control-plane   17d   v1.31.0
kubernetes-worker                  Ready    <none>         17d   v1.31.0
kubernetes-worker2                 Ready    <none>         17d   v1.31.0
manoj -->
manoj -->
manoj -->kubectl label node kubernetes-worker disktype=ssd
node/kubernetes-worker labeled
manoj -->
manoj -->kubectl get node --show-labels

NAME                                STATUS    ROLES    AGE   VERSION   LABELS
kubernetes-control-plane           Ready    control-plane   17d   v1.31.0   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/
osname=kubernetes-control-plane,kubernetes.io/os=linux,node-role.kubernetes.io/control-plane=node.kubernetes.io/exclude-from-external-load-balancers=
kubernetes-worker                  Ready    <none>         17d   v1.31.0   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,disktype=ssd,kubernetes.io/arch=amd64,k
ubernetes.io/hostname=kubernetes-worker,kubernetes.io/os=linux
kubernetes-worker2                 Ready    <none>         17d   v1.31.0   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/
osname=kubernetes-worker2,kubernetes.io/os=linux
manoj -->
manoj -->kubectl get pod -o wide

NAME    READY   STATUS    RESTARTS   AGE   IP          NODE            NOMINATED NODE   READINESS GATES
nginx   1/1     Running   0           10m   10.244.1.3  kubernetes-worker   <none>            <none>
manoj -->
manoj -->

```

```

/var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-7qvhv (ro)
Conditions:
  Type              Status
  PodReadyToStartContainers  True
  Initialized        True
  Ready              True
  ContainersReady     True
  PodScheduled       True

Volumes:
  kube-api-access-7qvhv:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
QoS Class:           BestEffort
Node-Selectors:      <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                     node.kubernetes.io/unreachable:NoExecute op=Exists for 300s

Events:
  Type      Reason                  Age    From                      Message
  ----      -
  Warning   FailedScheduling         12m    default-scheduler        0/3 nodes are available: 1 node(s) had untol
idn't match Pod's node affinity selector. preemption: 0/3 nodes are available: 3 Preemption is not helpful for scheduling.
  Warning   FailedScheduling         6m43s  default-scheduler        0/3 nodes are available: 1 node(s) had untol
idn't match Pod's node affinity selector. preemption: 0/3 nodes are available: 3 Preemption is not helpful for scheduling.
  Normal    Scheduled                2m4s   default-scheduler        Successfully assigned default/nginx to kubern
  Normal    Pulling                  2m4s   kubelet                  Pulling image "nginx"
  Normal    Pulled                   2m2s   kubelet                  Successfully pulled image "nginx" in 1.805s (1.805s including waiting). Image size: 72950394 bytes.
  Normal    Created                  2m2s   kubelet                  Created container nginx
  Normal    Started                  2m2s   kubelet                  Started container nginx

```

Note: we use node affinity, taint and toleration together to make sure pod are accumulating in the nodes that are meant for it.

Eg: if we have large node and that node is run's on particular type of workload like GPU specific workload or AIML specific workload or node with high performance that is only meant to run data warehousing workload etc, in those case we use this together.